

ZP_LC_pt1

October 11, 2018

1 Kaggle Lending Club Loan Data

Created: October 11, 2018

Author: Zack Pollak <zapollak@gmail.com>

1.1 Part 1

1.1.1 Import Python packages

```
In [112]: from __future__ import print_function
import datetime
import itertools
import re
import numpy as np
import pandas as pd
import sqlite3
import scipy
import sklearn
from sklearn import preprocessing as pp
import matplotlib.pyplot as plt
import seaborn as sns
from plotly.offline import download_plotlyjs, init_notebook_mode, plot, iplot
import plotly.graph_objs as go

# Ignore pandas warnings when writing slice back to orig df
pd.options.mode.chained_assignment = None # default='warn'
# Display all columns when viewing pandas df
pd.set_option('display.max_columns', None)
# Set plotly offline notebook mode
init_notebook_mode(connected=True)
```

1.1.2 Pull data into environment

```
In [3]: df = pd.read_csv('lc_loan_data/loan.csv', low_memory=False)
```

```
In [8]: df.head(2)
```

```

Out[8]:      id member_id loan_amnt funded_amnt funded_amnt_inv      term \
0 1077501    1296599    5000.0    5000.0    4975.0  36 months
1 1077430    1314167    2500.0    2500.0    2500.0  60 months

      int_rate installment grade sub_grade emp_title emp_length home_ownership \
0      10.65      162.87    B      B2      NaN  10+ years      RENT
1      15.27      59.83    C      C4    Ryder    < 1 year      RENT

      annual_inc verification_status issue_d loan_status pymnt_plan \
0      24000.0      Verified Dec-2011    Fully Paid      n
1      30000.0    Source Verified Dec-2011    Charged Off      n

      url \
0 https://www.lendingclub.com/browse/loanDetail...
1 https://www.lendingclub.com/browse/loanDetail...

      desc      purpose      title \
0 Borrower added on 12/22/11 > I need to upgra... credit_card Computer
1 Borrower added on 12/22/11 > I plan to use t...      car      bike

      zip_code addr_state      dti delinq_2yrs earliest_cr_line inq_last_6mths \
0      860xx      AZ  27.65      0.0      Jan-1985      1.0
1      309xx      GA  1.00      0.0      Apr-1999      5.0

      mths_since_last_delinq mths_since_last_record open_acc pub_rec \
0      NaN      NaN      3.0      0.0
1      NaN      NaN      3.0      0.0

      revol_bal revol_util total_acc initial_list_status out_prncp \
0      13648.0      83.7      9.0      f      0.0
1      1687.0      9.4      4.0      f      0.0

      out_prncp_inv total_pymnt total_pymnt_inv total_rec_prncp \
0      0.0  5861.071414      5831.78      5000.00
1      0.0  1008.710000      1008.71      456.46

      total_rec_int total_rec_late_fee recoveries collection_recovery_fee \
0      861.07      0.0      0.00      0.00
1      435.17      0.0      117.08      1.11

      last_pymnt_d last_pymnt_amnt next_pymnt_d last_credit_pull_d \
0      Jan-2015      171.62      NaN      Jan-2016
1      Apr-2013      119.66      NaN      Sep-2013

      collections_12_mths_ex_med mths_since_last_major_derog policy_code \
0      0.0      NaN      1.0
1      0.0      NaN      1.0

```

	application_type	annual_inc_joint	dti_joint	verification_status_joint	\
0	INDIVIDUAL	NaN	NaN	NaN	
1	INDIVIDUAL	NaN	NaN	NaN	

	acc_now_delinq	tot_coll_amt	tot_cur_bal	open_acc_6m	open_il_6m	\
0	0.0	NaN	NaN	NaN	NaN	
1	0.0	NaN	NaN	NaN	NaN	

	open_il_12m	open_il_24m	mths_since_rcnt_il	total_bal_il	il_util	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	

	open_rv_12m	open_rv_24m	max_bal_bc	all_util	total_rev_hi_lim	inq_fi	\
0	NaN	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	NaN	

	total_cu_tl	inq_last_12m
0	NaN	NaN
1	NaN	NaN

Let's check the dataframe dimensions.

```
In [26]: def df_shape(data):
          print('Number of columns: ' + str(data.shape[1]))
          print('Number of rows: ' + str(data.shape[0]))
```

```
In [27]: df_shape(df)
```

```
Number of columns: 74
Number of rows: 887379
```

```
In [13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 887379 entries, 0 to 887378
Data columns (total 74 columns):
id                887379 non-null int64
member_id         887379 non-null int64
loan_amnt         887379 non-null float64
funded_amnt       887379 non-null float64
funded_amnt_inv   887379 non-null float64
term              887379 non-null object
int_rate          887379 non-null float64
installment       887379 non-null float64
grade             887379 non-null object
sub_grade         887379 non-null object
emp_title         835922 non-null object
emp_length        887379 non-null object
```

home_ownership	887379	non-null	object
annual_inc	887375	non-null	float64
verification_status	887379	non-null	object
issue_d	887379	non-null	object
loan_status	887379	non-null	object
pymnt_plan	887379	non-null	object
url	887379	non-null	object
desc	126029	non-null	object
purpose	887379	non-null	object
title	887228	non-null	object
zip_code	887379	non-null	object
addr_state	887379	non-null	object
dti	887379	non-null	float64
delinq_2yrs	887350	non-null	float64
earliest_cr_line	887350	non-null	object
inq_last_6mths	887350	non-null	float64
mths_since_last_delinq	433067	non-null	float64
mths_since_last_record	137053	non-null	float64
open_acc	887350	non-null	float64
pub_rec	887350	non-null	float64
revol_bal	887379	non-null	float64
revol_util	886877	non-null	float64
total_acc	887350	non-null	float64
initial_list_status	887379	non-null	object
out_prncp	887379	non-null	float64
out_prncp_inv	887379	non-null	float64
total_pymnt	887379	non-null	float64
total_pymnt_inv	887379	non-null	float64
total_rec_prncp	887379	non-null	float64
total_rec_int	887379	non-null	float64
total_rec_late_fee	887379	non-null	float64
recoveries	887379	non-null	float64
collection_recovery_fee	887379	non-null	float64
last_pymnt_d	869720	non-null	object
last_pymnt_amnt	887379	non-null	float64
next_pymnt_d	634408	non-null	object
last_credit_pull_d	887326	non-null	object
collections_12_mths_ex_med	887234	non-null	float64
mths_since_last_major_derog	221703	non-null	float64
policy_code	887379	non-null	float64
application_type	887379	non-null	object
annual_inc_joint	511	non-null	float64
dti_joint	509	non-null	float64
verification_status_joint	511	non-null	object
acc_now_delinq	887350	non-null	float64
tot_coll_amt	817103	non-null	float64
tot_cur_bal	817103	non-null	float64
open_acc_6m	21372	non-null	float64

```

open_il_6m                21372 non-null float64
open_il_12m               21372 non-null float64
open_il_24m               21372 non-null float64
mths_since_rcnt_il        20810 non-null float64
total_bal_il              21372 non-null float64
il_util                   18617 non-null float64
open_rv_12m               21372 non-null float64
open_rv_24m               21372 non-null float64
max_bal_bc                21372 non-null float64
all_util                  21372 non-null float64
total_rev_hi_lim          817103 non-null float64
inq-fi                    21372 non-null float64
total_cu_tl               21372 non-null float64
inq_last_12m              21372 non-null float64
dtypes: float64(49), int64(2), object(23)
memory usage: 501.0+ MB

```

Looking at the initial information on the data shows that there are a lot of missing values. The rightmost part of the dataframe looks to be pretty sparse with a lot of fields having only 21,372 non-null values.

1.1.3 Subset columns for EDA

```

In [176]: fields = ['loan_amnt', 'funded_amnt', 'term', 'int_rate', 'grade', 'annual_inc',
                    'issue_d', 'dti', 'revol_bal', 'total_pymnt', 'loan_status']
df_eda = df[fields]

```

```

In [177]: df_eda.head()

```

```

Out[177]:   loan_amnt  funded_amnt      term  int_rate  grade  annual_inc  issue_d  \
0      5000.0      5000.0  36 months    10.65      B      24000.0  Dec-2011
1       2500.0      2500.0  60 months    15.27      C      30000.0  Dec-2011
2       2400.0      2400.0  36 months    15.96      C      12252.0  Dec-2011
3      10000.0     10000.0  36 months    13.49      C      49200.0  Dec-2011
4       3000.0       3000.0  60 months    12.69      B      80000.0  Dec-2011

          dti  revol_bal  total_pymnt  loan_status
0    27.65    13648.0    5861.071414  Fully Paid
1     1.00     1687.0    1008.710000  Charged Off
2     8.72     2956.0    3003.653644  Fully Paid
3    20.00     5598.0   12226.302212  Fully Paid
4    17.94    27783.0    3242.170000    Current

```

```

In [178]: df_shape(df_eda)

```

```

Number of columns: 11
Number of rows: 887379

```

1.1.4 Grab summary statistics for the new dataset

```
In [179]: df_eda.describe()
```

```
Out[179]:
```

	loan_amnt	funded_amnt	int_rate	annual_inc	\
count	887379.000000	887379.000000	887379.000000	8.873750e+05	
mean	14755.264605	14741.877625	13.246740	7.502759e+04	
std	8435.455601	8429.897657	4.381867	6.469830e+04	
min	500.000000	500.000000	5.320000	0.000000e+00	
25%	8000.000000	8000.000000	9.990000	4.500000e+04	
50%	13000.000000	13000.000000	12.990000	6.500000e+04	
75%	20000.000000	20000.000000	16.200000	9.000000e+04	
max	35000.000000	35000.000000	28.990000	9.500000e+06	

	dti	revol_bal	total_pymnt
count	887379.000000	8.873790e+05	887379.000000
mean	18.157039	1.692079e+04	7558.826684
std	17.190626	2.242679e+04	7871.243336
min	0.000000	0.000000e+00	0.000000
25%	11.910000	6.443000e+03	1914.590000
50%	17.650000	1.187500e+04	4894.999117
75%	23.950000	2.082900e+04	10616.814231
max	9999.000000	2.904836e+06	57777.579870

For all numerical variables, the mean is greater than the median indicating a positive skew.

1.1.5 Missing data

```
In [180]: print('Which columns have null values?\n')
          print(df_eda.isnull().any())

          print('\n')

          print('How many null values does each column have?\n')
          print(df_eda.isnull().sum().sort_values(ascending=False))
```

Which columns have null values?

loan_amnt	False
funded_amnt	False
term	False
int_rate	False
grade	False
annual_inc	True
issue_d	False
dti	False
revol_bal	False
total_pymnt	False
loan_status	False

```
dtype: bool
```

How many null values does each column have?

```
annual_inc      4
loan_status      0
total_pymnt      0
revol_bal        0
dti              0
issue_d          0
grade           0
int_rate        0
term            0
funded_amnt      0
loan_amnt        0
dtype: int64
```

There are only 4 missing values in the dataset. Let's take a look at them.

```
In [181]: df_eda[df_eda['annual_inc'].isnull()]
```

```
Out[181]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	\
42449	5000.0	5000.0	36 months	7.43	A	NaN	
42450	7000.0	7000.0	36 months	7.75	A	NaN	
42480	6700.0	6700.0	36 months	7.75	A	NaN	
42533	6500.0	6500.0	36 months	8.38	A	NaN	

	issue_d	dti	revol_bal	total_pymnt	\
42449	Aug-2007	1.0	0.0	5593.46	
42450	Aug-2007	1.0	0.0	7867.53	
42480	Jul-2007	1.0	0.0	7530.42	
42533	Jun-2007	4.0	0.0	7373.83	

	loan_status
42449	Does not meet the credit policy. Status:Fully ...
42450	Does not meet the credit policy. Status:Fully ...
42480	Does not meet the credit policy. Status:Fully ...
42533	Does not meet the credit policy. Status:Fully ...

All 4 rows with a missing value for **annual_inc** is a loan that does not meet the credit policy with notes that it was fully paid. Looking at the **total_pymnt** field, I would assume that these are old loans which have been paid, however, they no longer meet the current credit policy. As you need to provide your annual income to get a loan, I do not think that these are records where the borrower did not provide information. I think it is safe to remove these records from the dataset rather than perform imputation.

```
In [182]: init_len = df_eda.shape[0]
df_eda = df_eda.dropna()
```

```
new_len = df_eda.shape[0]
print('{} rows were dropped from the dataset. This is a {:.2f}% reduction of rows.'.format(100 * (old_len - new_len) / old_len))
```

4 rows were dropped from the dataset. This is a 0.00% reduction of rows.

```
In [183]: df_shape(df_eda)
```

Number of columns: 11

Number of rows: 887375

1.1.6 Distribution of categorical variables

We will take a look at the value counts for **grade**, **term**, and **loan_status**. While **issue_d** is currently strictly categorical due to the formatting. It makes sense to look at the distribution at the moment but I assume it will be dispersed across all months from 2007-2015. It is worth noting that it makes sense to transform **term** into a numerical variable by removing the 'months' suffix resulting in values of 36 or 60. It will clearly still only have 2 levels.

Grade (grade)

```
In [184]: print('Grade Counts')
pd.value_counts(df_eda['grade']).to_frame().reset_index()
```

Grade Counts

```
Out[184]:
```

index	grade	count
0	B	254535
1	C	245860
2	A	148198
3	D	139542
4	E	70705
5	F	23046
6	G	5489

```
In [185]: print('Grade Bucket Counts and Percentages\n\n')
print('Number of mid-grade loans (B/C/D): ' + str(df_eda[df_eda['grade'].isin(['B', 'C', 'D'])].shape[0]))
print('Percentage of total loans: {:.2f}%'.format(100 * df_eda[df_eda['grade'].isin(['B', 'C', 'D'])].shape[0] / df_eda.shape[0]))
print('\n')

print('Number of high-grade loans (A): ' + str(df_eda[df_eda['grade'] == 'A'].shape[0]))
print('Percentage of total loans: {:.2f}%'.format(100 * df_eda[df_eda['grade'] == 'A'].shape[0] / df_eda.shape[0]))
print('\n')

print('Number of low-grade loans (E/F/G): ' + str(df_eda[df_eda['grade'].isin(['E', 'F', 'G'])].shape[0]))
print('Percentage of total loans: {:.2f}%'.format(100 * df_eda[df_eda['grade'].isin(['E', 'F', 'G'])].shape[0] / df_eda.shape[0]))
```


Grade Bucket Counts and Percentages

Number of mid-grade loans (B/C/D): 639937
Percentage of total loans: 72.12%

Number of high-grade loans (A): 148198
Percentage of total loans: 16.70%

Number of low-grade loans (E/F/G): 99240
Percentage of total loans: 11.18%

A majority of loans are mid-grade (B/C/D), followed by high-grade (A) loans with risky, low-grade (E/F/G) loans making up the tail. This distribution makes sense as I would not expect a majority of loans to be high-grade and at the same time Lending Club does not want a lot of risky loans.

Term (term)

```
In [186]: print('Term Counts')
          pd.value_counts(df_eda['term']).to_frame().reset_index()
```

Term Counts

```
Out[186]:
```

	index	term
0	36 months	621121
1	60 months	266254

```
In [187]: print('Term Percentages')
          pd.value_counts(df_eda['term'], normalize=True).to_frame().reset_index()
```

Term Percentages

```
Out[187]:
```

	index	term
0	36 months	0.699953
1	60 months	0.300047

The loans in the dataset only have 36-month and 60-month terms as noted in the data dictionary. 70% of the loans have 36-month terms while the remaining 30% have 60-month terms. This is a clear indication that a majority of the loans through Lending Club are shorter-term loans.

Loan status (loan_status)

```
In [188]: print('Loan Status Counts')
          pd.value_counts(df_eda['loan_status']).to_frame().reset_index()
```

Loan Status Counts

```
Out[188]:
```

	index	loan_status
0	Current	601779
1	Fully Paid	207723
2	Charged Off	45248
3	Late (31-120 days)	11591
4	Issued	8460
5	In Grace Period	6253
6	Late (16-30 days)	2357
7	Does not meet the credit policy. Status:Fully ...	1984
8	Default	1219
9	Does not meet the credit policy. Status:Charge...	761

```
In [189]: print('Loan Status Percentages')
          pd.value_counts(df_eda['loan_status'], normalize=True).to_frame().reset_index()
```

Loan Status Percentages

```
Out[189]:
```

	index	loan_status
0	Current	0.678156
1	Fully Paid	0.234087
2	Charged Off	0.050991
3	Late (31-120 days)	0.013062
4	Issued	0.009534
5	In Grace Period	0.007047
6	Late (16-30 days)	0.002656
7	Does not meet the credit policy. Status:Fully ...	0.002236
8	Default	0.001374
9	Does not meet the credit policy. Status:Charge...	0.000858

A majority of the loans are current (67.8%) with the second largest group being those that have been fully paid (23.4%). There are a fair amount of loans which have been charged off (5%) and an even smaller percentage which are late (~1.5%). Loans that have defaulted constitute a very small portion of the dataset (0.1%) which is a good sign for Lending Club. There are still loans in the dataset with a loan status indicating that they do not meet the credit policy (2,745 or 0.3%). From a consistency standpoint, it would make sense to remove these loans as we did earlier with the loans that were missing data for **annual_inc**. ##### Let's remove the loans that do not meet the credit policy

```
In [190]: init_len = df_eda.shape[0]
          df_eda = df_eda[~df_eda['loan_status'].str.contains('Does not meet', na=False)]
          new_len = df_eda.shape[0]
          print('{} rows were dropped from the dataset. This is a {:.2f}% reduction of rows.'.format(
```

2745 rows were dropped from the dataset. This is a 0.31% reduction of rows.

Issue Date (issue_d)

```
In [191]: print('Issue Date counts')
          pd.value_counts(df_eda['issue_d']).to_frame().reset_index()
```

Issue Date counts

```
Out[191]:
```

	index	issue_d
0	Oct-2015	48631
1	Jul-2015	45962
2	Dec-2015	44342
3	Oct-2014	38782
4	Nov-2015	37530
5	Aug-2015	35886
6	Apr-2015	35427
7	Jan-2015	35107
8	May-2015	31913
9	Jul-2014	29306
10	Sep-2015	28641
11	Jun-2015	28485
12	Mar-2015	25400
13	Nov-2014	25054
14	Feb-2015	23770
15	May-2014	19099
16	Apr-2014	19071
17	Aug-2014	18814
18	Jun-2014	17179
19	Mar-2014	16513
20	Jan-2014	15628
21	Feb-2014	15269
22	Dec-2013	15020
23	Nov-2013	14676
24	Oct-2013	14114
25	Sep-2013	12987
26	Aug-2013	12674
27	Jul-2013	11910
28	Jun-2013	10899
29	Sep-2014	10606
..
73	Jan-2010	589
74	Oct-2009	545
75	Sep-2009	449
76	Aug-2009	408
77	Jul-2009	374

78	Jun-2009	356
79	May-2009	319
80	Apr-2009	290
81	Mar-2009	276
82	Feb-2009	260
83	Jan-2009	239
84	Mar-2008	236
85	Dec-2008	223
86	Nov-2008	184
87	Feb-2008	174
88	Jan-2008	171
89	Apr-2008	155
90	Oct-2008	96
91	Dec-2007	85
92	Jul-2008	83
93	May-2008	71
94	Aug-2008	71
95	Jun-2008	66
96	Oct-2007	47
97	Nov-2007	37
98	Aug-2007	33
99	Sep-2008	32
100	Jul-2007	30
101	Sep-2007	18
102	Jun-2007	1

[103 rows x 2 columns]

```
In [192]: print('Issue Date percentages')
pd.value_counts(df_eda['issue_d'], normalize=True).to_frame().reset_index()
```

Issue Date percentages

```
Out[192]:
```

	index	issue_d
0	Oct-2015	0.054973
1	Jul-2015	0.051956
2	Dec-2015	0.050125
3	Oct-2014	0.043840
4	Nov-2015	0.042425
5	Aug-2015	0.040566
6	Apr-2015	0.040047
7	Jan-2015	0.039686
8	May-2015	0.036075
9	Jul-2014	0.033128
10	Sep-2015	0.032376
11	Jun-2015	0.032200
12	Mar-2015	0.028713

13	Nov-2014	0.028321
14	Feb-2015	0.026870
15	May-2014	0.021590
16	Apr-2014	0.021558
17	Aug-2014	0.021268
18	Jun-2014	0.019419
19	Mar-2014	0.018667
20	Jan-2014	0.017666
21	Feb-2014	0.017260
22	Dec-2013	0.016979
23	Nov-2013	0.016590
24	Oct-2013	0.015955
25	Sep-2013	0.014681
26	Aug-2013	0.014327
27	Jul-2013	0.013463
28	Jun-2013	0.012320
29	Sep-2014	0.011989
..
73	Jan-2010	0.000666
74	Oct-2009	0.000616
75	Sep-2009	0.000508
76	Aug-2009	0.000461
77	Jul-2009	0.000423
78	Jun-2009	0.000402
79	May-2009	0.000361
80	Apr-2009	0.000328
81	Mar-2009	0.000312
82	Feb-2009	0.000294
83	Jan-2009	0.000270
84	Mar-2008	0.000267
85	Dec-2008	0.000252
86	Nov-2008	0.000208
87	Feb-2008	0.000197
88	Jan-2008	0.000193
89	Apr-2008	0.000175
90	Oct-2008	0.000109
91	Dec-2007	0.000096
92	Jul-2008	0.000094
93	May-2008	0.000080
94	Aug-2008	0.000080
95	Jun-2008	0.000075
96	Oct-2007	0.000053
97	Nov-2007	0.000042
98	Aug-2007	0.000037
99	Sep-2008	0.000036
100	Jul-2007	0.000034
101	Sep-2007	0.000020
102	Jun-2007	0.000001

```
[103 rows x 2 columns]
```

As suspected, the loans are distributed across all months from mid-2007 through the end of 2015. Interestingly, the latter years tend to contain a majority of the loans. I assume this is due to the financial crisis of 2007-08.

1.1.7 Transformations

Let's transform the non-numerical fields from **df_eda** into numerical fields so that we may run some more in-depth analysis.

Remove the 'months' string from **term** (e.g. '36 months' to 36)

Encode **grade** to numerical (e.g. 'A' = 0, 'B' = 1, etc.)

Transform **issue_d** into date

Encode **loan_status** as numerical (e.g. 'Fully Paid' = 0, 'Charged Off' = 1, etc.)

```
In [197]: df_eda.head()
```

```
Out[197]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	issue_d	\
0	5000.0	5000.0	36 months	10.65	B	24000.0	Dec-2011	
1	2500.0	2500.0	60 months	15.27	C	30000.0	Dec-2011	
2	2400.0	2400.0	36 months	15.96	C	12252.0	Dec-2011	
3	10000.0	10000.0	36 months	13.49	C	49200.0	Dec-2011	
4	3000.0	3000.0	60 months	12.69	B	80000.0	Dec-2011	

	dti	revol_bal	total_pymnt	loan_status
0	27.65	13648.0	5861.071414	Fully Paid
1	1.00	1687.0	1008.710000	Charged Off
2	8.72	2956.0	3003.653644	Fully Paid
3	20.00	5598.0	12226.302212	Fully Paid
4	17.94	27783.0	3242.170000	Current

```
In [198]: df_num = df_eda.copy()
df_num.head()
```

```
Out[198]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	issue_d	\
0	5000.0	5000.0	36 months	10.65	B	24000.0	Dec-2011	
1	2500.0	2500.0	60 months	15.27	C	30000.0	Dec-2011	
2	2400.0	2400.0	36 months	15.96	C	12252.0	Dec-2011	
3	10000.0	10000.0	36 months	13.49	C	49200.0	Dec-2011	
4	3000.0	3000.0	60 months	12.69	B	80000.0	Dec-2011	

	dti	revol_bal	total_pymnt	loan_status
0	27.65	13648.0	5861.071414	Fully Paid
1	1.00	1687.0	1008.710000	Charged Off
2	8.72	2956.0	3003.653644	Fully Paid
3	20.00	5598.0	12226.302212	Fully Paid
4	17.94	27783.0	3242.170000	Current

```

In [199]: # term
df_num['term'] = df_num['term'].apply(lambda x: re.findall(r'\d+', x))
df_num['term'] = df_num['term'].apply(lambda x: int(x[0]))

# grade
le = pp.LabelEncoder()
le.fit(df_num['grade'])
df_num['grade'] = le.transform(df_num['grade'])

# issue_d
df_num['issue_d'] = pd.to_datetime(df_num['issue_d'])

# loan_status
le = pp.LabelEncoder()
le.fit(df_num['loan_status'])
df_num['loan_status'] = le.transform(df_num['loan_status'])

In [200]: df_num.head()

```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	issue_d	\
0	5000.0	5000.0	36	10.65	1	24000.0	2011-12-01	
1	2500.0	2500.0	60	15.27	2	30000.0	2011-12-01	
2	2400.0	2400.0	36	15.96	2	12252.0	2011-12-01	
3	10000.0	10000.0	36	13.49	2	49200.0	2011-12-01	
4	3000.0	3000.0	60	12.69	1	80000.0	2011-12-01	

	dti	revol_bal	total_pymnt	loan_status
0	27.65	13648.0	5861.071414	3
1	1.00	1687.0	1008.710000	0
2	8.72	2956.0	3003.653644	3
3	20.00	5598.0	12226.302212	3
4	17.94	27783.0	3242.170000	1


1.1.8 Correlation matrix (heatmap)

```

In [201]: %matplotlib inline
plt.figure(figsize=(20,10))
sns.set_context("notebook", font_scale=1.5)
sns.heatmap(df_num.corr(), vmax=1, square=True, annot=True, fmt='.2f')

Out[201]: <matplotlib.axes._subplots.AxesSubplot at 0x15b2fce90>

```



Summary: **loan_amnt** and **funded_amnt** have a perfect correlation of 1

The second highest correlation is between **grade** and **int_rate** (0.96); this more than likely implies that one is used to determine the other

grade is also decently correlated with **term** (0.45); regarding the last bullet, **int_rate** is also rather correlated with **term** though not identical to **grade/term** correlation (0.43)


There is pretty strong correlation between **total_pymnt** and **loan_amnt/funded_amnt** (0.47)

There aren't many variables that are inversely correlated with the largest being **dti** and **annual_inc** (-0.09)

1.1.9 Average Loan Amount by Grade

```
In [213]: by_grade = df_eda.groupby('grade').mean()
          avg_loan_amnt = by_grade['loan_amnt'].reset_index()
          plt.subplots(figsize=(15,10))
          sns.barplot(x='grade', y='loan_amnt', data=avg_loan_amnt)
```

Out [213]: <matplotlib.axes._subplots.AxesSubplot at 0x190ed2590>



```
In [214]: avg_loan_amnt.sort_values('loan_amnt', ascending=False)
```

```
Out[214]:
```

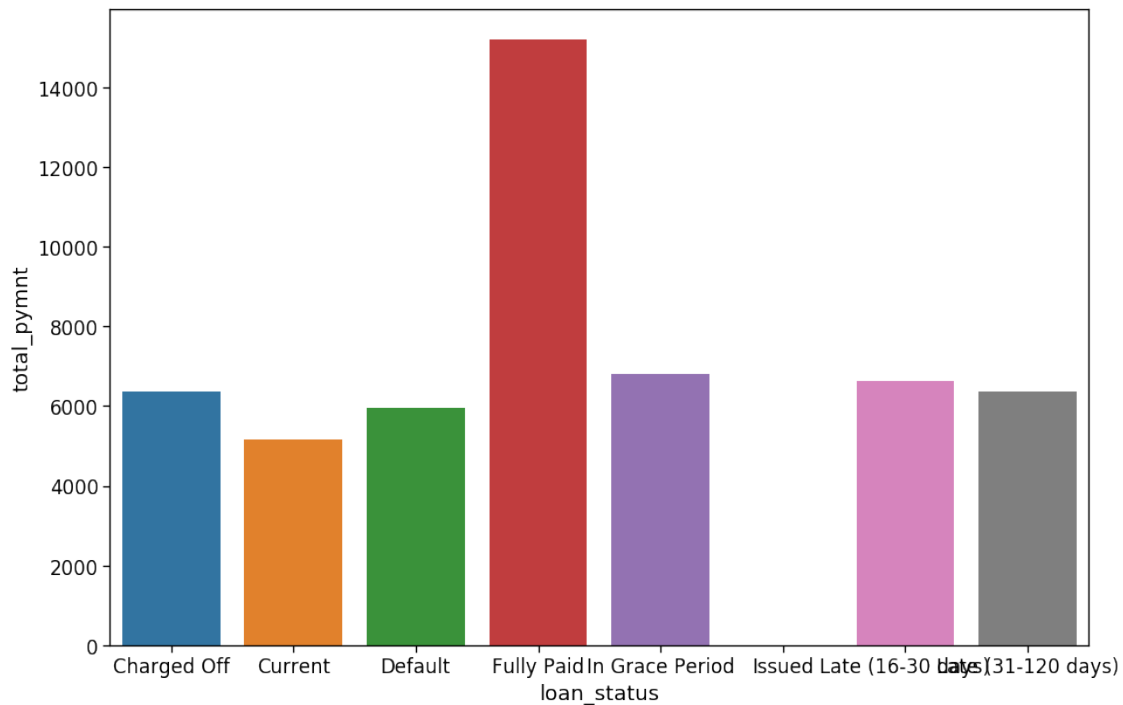
	grade	loan_amnt
6	G	20855.868744
5	F	19217.782578
4	E	18071.066995
3	D	15494.357081
2	C	14480.310503
0	A	14044.276151
1	B	13650.167695

Summary When looking at loan amounts (**loan_amnt**) by **grade**, we can see that riskier loans have higher average loan amounts. This seems like a recipe for disaster but at the same time, I am led to believe that **loan_amnt** is used to calculate the **grade** and higher loan amounts will result in a lower **grade**. The only outlier to this trend is that the average loan amount for grade A loans is larger than grade B loans.

1.1.10 Average Total Payment by Loan Status

```
In [218]: by_status = df_eda.groupby('loan_status').mean()
          avg_total_pymnt = by_status['total_pymnt'].reset_index()
          plt.subplots(figsize=(15,10))
          sns.barplot(x='loan_status', y='total_pymnt', data=avg_total_pymnt)
```

```
Out[218]: <matplotlib.axes._subplots.AxesSubplot at 0x191538d50>
```



```
In [220]: avg_total_pymnt.sort_values('total_pymnt', ascending=False)
```

```
Out[220]:
```


	loan_status	total_pymnt
3	Fully Paid	15186.680773
4	In Grace Period	6792.092241
6	Late (16-30 days)	6617.883309
0	Charged Off	6369.612648
7	Late (31-120 days)	6367.870714
2	Default	5947.982329
1	Current	5152.886189
5	Issued	9.307895

Summary We dug into the total payment amount (**total_pymnt**) by **loan_status** because **total_pymnt** had the highest correlation with **loan_status** out of all variables. One can see that the Fully Paid loans have the largest average total payment. This makes perfect sense as the loans were paid off. The loans with a status indicating they are late have a higher average total payment amount than those with a status of 'Current' which is probably due to interest.

1.1.11 Distribution of Loan Amount

```
In [221]: plt.subplots(figsize=(15,10))
          sns.distplot(df['loan_amnt'])
```

Out[221]: <matplotlib.axes._subplots.AxesSubplot at 0x1918433d0>



```
In [227]: from scipy.stats import skew, kurtosis
          print('The skewness of 'loan_amnt' is: {:.3f}'.format(skew(df_eda['loan_amnt'])))
          print('The excess kurtosis of 'loan_amnt' is: {:.3f}'.format(kurtosis(df_eda['loan_a
```

The skewness of loan_amnt is: 0.681

The excess kurtosis of loan_amnt is: -0.259

Summary As seen in the **loan_amnt** distribution plot, there is tail to the right; the positive skew value of 0.681 confirms this

The kurtosis value of -0.259 indicates that the peakedness is pretty close to normal

1.2 Part 2

Create a dataset containing strictly 36 month loans.

```
In [230]: df_36 = df_eda.copy()
          df_36 = df_36[df_36['term'].str.contains('36')]
          df_36_num = df_num.copy()
          df_36_num = df_36_num[df_36_num['term'] == 36]
```

```
In [235]: print(df_36.shape)
          df_36.head()
```

(618687, 11)

```
Out[235]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	issue_d	\
0	5000.0	5000.0	36 months	10.65	B	24000.0	Dec-2011	
2	2400.0	2400.0	36 months	15.96	C	12252.0	Dec-2011	
3	10000.0	10000.0	36 months	13.49	C	49200.0	Dec-2011	
5	5000.0	5000.0	36 months	7.90	A	36000.0	Dec-2011	
7	3000.0	3000.0	36 months	18.64	E	48000.0	Dec-2011	

	dti	revol_bal	total_pymnt	loan_status
0	27.65	13648.0	5861.071414	Fully Paid
2	8.72	2956.0	3003.653644	Fully Paid
3	20.00	5598.0	12226.302212	Fully Paid
5	11.20	7963.0	5631.377753	Fully Paid
7	5.35	8221.0	3938.144334	Fully Paid

```
In [236]: print('Term Counts')
pd.value_counts(df_36['term']).to_frame().reset_index()
```

```
Out[236]:
```

	index	term
0	36 months	618687

```
In [234]: print(df_36_num.shape)
df_36_num.head()
```

(618687, 11)

```
Out[234]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	issue_d	\
0	5000.0	5000.0	36	10.65	1	24000.0	2011-12-01	
2	2400.0	2400.0	36	15.96	2	12252.0	2011-12-01	
3	10000.0	10000.0	36	13.49	2	49200.0	2011-12-01	
5	5000.0	5000.0	36	7.90	0	36000.0	2011-12-01	
7	3000.0	3000.0	36	18.64	4	48000.0	2011-12-01	

	dti	revol_bal	total_pymnt	loan_status
0	27.65	13648.0	5861.071414	3
2	8.72	2956.0	3003.653644	3
3	20.00	5598.0	12226.302212	3
5	11.20	7963.0	5631.377753	3
7	5.35	8221.0	3938.144334	3

```
In [237]: print('Term Counts')
pd.value_counts(df_36_num['term']).to_frame().reset_index()
```

Term Counts

```
Out[237]:
```

	index	term
0	36	618687

After performing sanity check, data looks good to proceed.

1.2.1 #1

Percentage of loans fully paid.

```
In [240]: print('Number of Fully Paid 36-month loans: ' + str(df_36[df_36['loan_status'] == 'Fully Paid']))
          print('Percentage of total loans: {:.2f}%'.format(100*df_36[df_36['loan_status'] == 'Fully Paid'].count()/df_36.count()))
```

Number of Fully Paid 36-month loans: 167575

Percentage of total loans: 27.09%

27.09% of the 36-month loans have been fully paid.

1.2.2 #2

Bucket by origination year and grade. Which bucket has highest rate of default? Assume that any loan not fully paid has defaulted.

```
In [248]: # Convert to date then year
df_36['orig_year'] = pd.to_datetime(df_num['issue_d'])
df_36['orig_year'] = df_36['orig_year'].apply(lambda x: x.year)

# Group year and grade
df_36['year_grade'] = df_36['orig_year'].map(str) + ' ' + df_36['grade'].map(str)

# Create default variable; if 'Fully Paid' then 'No' else 'Yes'
df_36['default'] = np.where(df_36['loan_status'] == 'Fully Paid', 'No', 'Yes')
```

```
In [269]: df_36_default = df_36[df_36['default'] == 'Yes']
pd.value_counts(df_36_default['default']).to_frame().reset_index()
```

```
Out[269]:   index  default
0    Yes    451112
```

```
In [271]: pd.value_counts(df_36_default['year_grade'], normalize=True).to_frame().reset_index()
```

```
Out[271]:   index  year_grade
0    2015 B    0.193065
1    2015 C    0.160016
2    2015 A    0.148515
3    2014 B    0.086963
4    2014 C    0.072113
5    2015 D    0.066977
6    2014 A    0.059129
7    2013 B    0.047252
8    2014 D    0.033814
9    2013 C    0.029966
```

Loans that originated in 2015 with a B grade constitute the largest portion of loans that have defaulted.

1.2.3 #3

Bucket by origination year and grade. What annualized rate of return have the loans generated on average? Assume that Annualized Rate of Return = $(\text{total_pymnt} / \text{funded_amnt} - 1)^{(1/3)}$.

```
In [283]: from __future__ import division
          df_36['ann_rate'] = (df_36['total_pymnt'] / df_36['funded_amnt'] - 1) ** (1/3)

In [292]: print('Overall average rate of return: ' + str(df_36['ann_rate'].mean()))

          by_year_grade = df_36.groupby('year_grade').mean()
          avg_ann_rate = by_year_grade['ann_rate'].reset_index()
          avg_ann_rate.sort_values('ann_rate', ascending=False).head(10)
```

Overall average rate of return: 0.475485541756

```
Out[292]:
```

	year_grade	ann_rate
33	2011 G	0.683785
40	2012 G	0.674487
26	2010 G	0.671978
39	2012 F	0.661811
38	2012 E	0.648165
32	2011 F	0.628126
31	2011 E	0.627589
11	2008 F	0.618044
37	2012 D	0.617703
12	2008 G	0.612983

Looking at all of the loans, the average annual rate of return was 0.475.

Looking at buckets, loans that originated in 2011 (0.684), 2012 (0.674), and 2010 (0.672) with a G grade had the highest average annual rates of return.

```
In [ ]:
```